


# <https://github.com/je-clark/DFC19-PacketAnalysis>

## Example 1: Split a packet capture based on time

Take a large packet capture with known timestamps of interest, and break those times out for manual analysis

### Required knowledge: PCAP File Format

Data Sourced from <https://wiki.wireshark.org/Development/LibpcapFileFormat>



### Global Header Format

```
typedef struct pcap_hdr_s {
    guint32 magic_number;    /* magic number */
    guint16 version_major;  /* major version number */
    guint16 version_minor;  /* minor version number */
    gint32  thiszone;        /* GMT to local correction */
    guint32 sigfigs;         /* accuracy of timestamps */
    guint32 snaplen;         /* max length of captured packets, in
    guint32 network;         /* data link type */
} pcap_hdr_t;
```

### Packet Header Format

```
typedef struct pcaprec_hdr_s {
    guint32 ts_sec;           /* timestamp seconds */
    guint32 ts_usec;         /* timestamp microseconds */
    guint32 incl_len;        /* number of octets of packet saved in file */
    guint32 orig_len;        /* actual length of packet */
} pcaprec_hdr_t;
```

## What Do We Care About?

1. Magic Number: *Indicates the endianness of the system that generated the PCAP*
2. Timestamp (seconds): *Compare this with the given timestamp of interest*
3. Included Length: *Necessary to move the packet contents to a new file*

## Let's Do It!

```
[ ] # Imports
import sys, os, argparse
from datetime import datetime
```

```
[ ] def cli_args():
    ep= '''\
Default output file is <infile>_<anchor_time>_<duration>. To
overwrite, use the -o flag.

This will use local system time to compute timestamps. If the
capture originated in
a different time zone, convert the anchor_time argument to local
time when calling this application.
'''
    desc = 'Automatically slice PCAP files by time.'
    parser = argparse.ArgumentParser(description=desc, epilog=ep,
formatter_class=argparse.RawTextHelpFormatter)
    parser.add_argument('infile', action='store', help='Input
file. Must be a PCAP')
    parser.add_argument('anchor_time', nargs=2, action='store',
help='Start time or center time. Use format mm/dd/yy hh:mm:ss')
    parser.add_argument('duration', action='store',
help='Duration to end time, or duration from center time when
using -s. In seconds')
    parser.add_argument('-s', '--symmetric', action='store_true',
help='Interpret the anchor time as the center point')
    parser.add_argument('-o', '--outfile', action='store',
help='Override the default output file name')

    return parser.parse_args()
```

```
[ ] def determine_endianness(global_header): # Expect this to be a
byte array
    magic_num_test = b'\xa1\xb2\xc3\xd4'
    magic_num = global_header[0:4]
    if magic_num == magic_num_test:
        return 'big'
    else:
        return 'little'
```

```
[ ] def pcap_split(infile, outfile, anchor_time, duration,
symmetric=False): # Expecting most of this to be command line
input, so strings
    # Get start and end time
    duration = int(duration)
    anchor = int(datetime.strptime(' '.join(anchor_time),
r'%m/%d/%y %H:%M:%S').timestamp())
    if symmetric:
        start_time = anchor - duration
        end_time = anchor + duration
    else:
        start_time = anchor
        end_time = anchor + duration

    print(f'This uses UTC timestamps for start time
({start_time}) and end time ({end_time})')

    # Open input and output files
    in_ptr = open(infile, 'rb+')
    out_ptr = open(outfile, 'wb+')

    # Global Header
    global_header = in_ptr.read(24)
    endianness = determine_endianness(global_header)
    print(f'This file was generated on a {endianness}-endian
computer')
    out_ptr.write(global_header)

    # Packets
    while True:
        # Packet Header
        packet_header = in_ptr.read(16)
        if not packet_header:
            break # Python returns '' for EOF, which evaluates
FALSE

        timestamp = int.from_bytes(packet_header[0:4],
byteorder=endianness)
        packet_length = int.from_bytes(packet_header[8:12],
byteorder=endianness)
```

```

        # Determine if we care about this packet
        if timestamp > end_time: break

    if start_time <= timestamp:
        packet = in_ptr.read(packet_length)
        out_ptr.write(packet_header)
        out_ptr.write(packet)
    else:
        # Move the file pointer to the beginning of the
        # next packet header
        in_ptr.seek(packet_length,1)

    # Be nice to your pointers. Put them away when you're done
    with them
    in_ptr.close()
    out_ptr.close()

    return

```

## Let's start messing with it

My sample capture is './sample.pcap'

Starting date/time is Apr 18, 2019 9:33:03

Ending date/time is Apr 26, 2019 11:25:07

```

[ ] # Just get everything from April 19?
pcap_split('./sample.pcap', 'apr_19.pcap',
['04/19/19','00:00:00'], '86400')

```

```

[ ] # Let's get April 20 symmetrically
pcap_split('./sample.pcap', 'apr_20_symmetric.pcap',
['04/20/19','12:00:00'], '43200', symmetric=True)

```

## Example 2: TLS Version Audit

**Audit the TLS version used by incoming connections to the server so we can upgrade/disable insecure protocols without relying on a scream test**

# Required Knowledge: Packet Layer Headers

## Ethernet Header

Modern computers use Type II Ethernet




Image from [https://en.wikipedia.org/wiki/Ethernet\\_frame](https://en.wikipedia.org/wiki/Ethernet_frame)

## IP Header

We're only going to work with IPv4 for this example. It's complicated enough already.

IPv4 Header Format																																	
Offsets	Octet	0								1								2								3							
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	0	Version				IHL				DSCP						ECN		Total Length															
4	32	Identification																Flags				Fragment Offset											
8	64	Time To Live								Protocol								Header Checksum															
12	96	Source IP Address																															
16	128	Destination IP Address																															
20	160	Options (if IHL > 5)																															
24	192																																
28	224																																
32	256																																

Image screenshoted from <https://en.wikipedia.org/wiki/IPv4>

## TCP Header

There's still more.

TCP Header																																	
Offsets	Octet	0								1								2								3							
Octet	Bit	7	6	5	4	3	2	1	0	7	6	5	4	3	2	1	0	7	6	5	4	3	2	1	0	7	6	5	4	3	2	1	0
0	0	Source port																Destination port															
4	32	Sequence number																															
8	64	Acknowledgment number (if ACK set)																															
12	96	Data offset				Reserved 000			N S	C W R	E C E	U R G	A C K	P S H	R S T	S Y N	F I N	Window Size															
16	128	Checksum																Urgent pointer (if URG set)															
20	160	Options (if data offset > 5. Padded at the end with "0" bytes if necessary.)																															
...	...	...																															

Image screenshoted from [https://en.wikipedia.org/wiki/Transmission\\_Control\\_Protocol](https://en.wikipedia.org/wiki/Transmission_Control_Protocol)

## TLS Record Layer

There are many components to TLS, so the record layer keeps track of them.

Byte	0	1	2	3	4
0	Content Type	Version		Length	
1	Payload				

## TLS Handshake Layer

These are just the parts we care about for this application.

Byte	0	1	2	3	4	5
0	Handshake Type	Length			Version	

## General Approach

- The final, negotiated TLS version is reported in the Server Hello. We need to find those and extract the client's information
- Emphasize speed of execution. Interpret values out of order to aggressively skip packets

## Let's Go!

```
[ ] # Set up classes as a useful structure to store information and
    # retrieve it later

class TLSSummary:
    def __init__(self, src_ip, dst_ip, src_port, dst_port,
tls_version):
        self.src_ip = src_ip
        self.dst_ip = dst_ip
        self.src_port = src_port
        self.dst_port = dst_port
        self.tls_version = tls_version

    def get_src_ip(self):
        return self.src_ip

    def get_dst_ip(self):
        return self.dst_ip

    def get_src_port(self):
        return self.src_port

    def get_dst_port(self):
```

```

        return self.dst_port

    def get_tls_version(self):
        return self.tls_version

    def print_summary(self):
        sum = '''
        Source IP:      {}
        Destination IP: {}
        Source Port:     {}
        Destination Port: {}
        TLS Version:     {}
        '''.format(self.src_ip, self.dst_ip, self.src_port,
self.dst_port, self.tls_version)
        print(sum)

```

```
[ ] # A few helper functions
```

```

def bytes_to_ip_str(ip_in_bytes):

    first_octet = str(int.from_bytes(ip_in_bytes[0:1], byteorder
= 'big'))
    second_octet = str(int.from_bytes(ip_in_bytes[1:2], byteorder
= 'big'))
    third_octet = str(int.from_bytes(ip_in_bytes[2:3], byteorder
= 'big'))
    fourth_octet = str(int.from_bytes(ip_in_bytes[3:4], byteorder
= 'big'))

    return '.'.join([first_octet, second_octet, third_octet,
fourth_octet])

def is_server_hello(packet):

    # We know we have some things between us and the TLS layer
    # Let's deal with them one at a time

    # Ethernet Header - Always 14 bytes, unless we have VLANs.
    That's beyond the scope of this script.
    eth_hdr_len = 14

    # IP Header - Length defined in the IHL field in words.
    WORDS_TO_BYTES = 4 # Conversion factor
    IHL_MASK = int('00001111',2) # IHL field isn't a full byte,
so we need to use a bit mask to get the value
    ver_ihl_field = packet[eth_hdr_len:eth_hdr_len + 1] #
Version/IHL field is the first one in the IP Header
    ihl = int.from_bytes(ver_ihl_field, byteorder = 'big') &
IHL_MASK # Bitwise AND
    ip_hdr_len = ihl * WORDS_TO_BYTES

```

```

# TCP Header - Length defined in data offset field
DATA_OFFSET_MAST = int('11110000',2)
data_offset_byte = packet[eth_hdr_len + ip_hdr_len + 12:
eth_hdr_len + ip_hdr_len + 12 + 1]
# Since the data offset is upper 4 bits of the byte, we need
to shift it down to get the right value
tcp_hdr_len = ((int.from_bytes(data_offset_byte, byteorder =
'big') & DATA_OFFSET_MAST) >> 4) * WORDS_TO_BYTES

# Look at TLS Layer
tls_layer = eth_hdr_len + ip_hdr_len + tcp_hdr_len
TLS_HANDSHAKE = b'\x16' # Evaluates to an integer value of 22
SERVER_HELLO = b'\x02'
tls_content_type = packet[tls_layer:tls_layer+1]
if tls_content_type == TLS_HANDSHAKE:
    handshake_type = packet[tls_layer + 5: tls_layer + 5 + 1]
    if handshake_type == SERVER_HELLO:
        tls_version = packet[tls_layer + 9: tls_layer + 9 +
2]

        return (True, tls_version, tcp_hdr_len, ip_hdr_len)
        # So yes, this violates "do one thing well", but I'm
already here, already have the info,
        # and I've already told you I'm bad at coding
    return (False, '',0,0) # It's not what I wanted, so I'm mad
at the world and don't wanna give info

def get_ip_addrs(ip_hdr):
    src_ip_bytes = ip_hdr[12:16]
    dst_ip_bytes = ip_hdr[16:20]

    src_ip = bytes_to_ip_str(src_ip_bytes)
    dst_ip = bytes_to_ip_str(dst_ip_bytes)

    return src_ip, dst_ip

def get_tcp_ports(tcp_hdr):
    src_port_bytes = tcp_hdr[0:2]
    dst_port_bytes = tcp_hdr[2:4]

    src_port = int.from_bytes(src_port_bytes, byteorder='big')
    dst_port = int.from_bytes(dst_port_bytes, byteorder='big')

    return src_port, dst_port

def readable_tls_version(tls_version_bytes):

    if tls_version_bytes == b'\x03\x01':
        return "TLS 1.0"
    elif tls_version_bytes == b'\x03\x02':
        return "TLS 1.1"
    elif tls_version_bytes == b'\x03\x03':
        return "TLS 1.2"

```



```

elif tls_version_bytes == b'\x03\x00':
    return "SSL 3.0"
elif tls_version_bytes == b'\x03\x04':
    return "TLS 1.3"
else:
    print(f'Unknown version displayed as
{tls_version_bytes.hex()}')
    return "Unknown"

```

```

[ ] def tls_version_audit(pcap_file):

    # Open File
    pcap_ptr = open(pcap_file, 'rb+')

    # Get Endianness
    endianness = determine_endianness(pcap_ptr.read(24))

    # Deal with the packets
    tls_connection_list = []
    while True:

        # Get a packet
        pkt_hdr = pcap_ptr.read(16)
        if not pkt_hdr:
            break
        pkt_len = int.from_bytes(pkt_hdr[8:12],
byteorder=endianness)
        pkt = pcap_ptr.read(pkt_len)

        # Is it a Server Hello?
        resp = is_server_hello(pkt)
        if not resp[0]:
            continue # Go to the next packet
        # I would wrap this in an else:, but I don't need to
        tls_version_bytes = resp[1]
        tcp_hdr_len = resp[2]
        ip_hdr_len = resp[3]

        # Gather necessary values
        # Really only need server port and client IP, but
collecting both source and destination because I can
        ip_hdr = pkt[14: 14 + ip_hdr_len]
        tcp_hdr = pkt[14 + ip_hdr_len: 14 + ip_hdr_len +
tcp_hdr_len]

        src_ip, dst_ip = get_ip_addrs(ip_hdr)
        src_port, dst_port = get_tcp_ports(tcp_hdr)
        tls_version = readable_tls_version(tls_version_bytes)

```

```
        tls_connection_list.append(TLSSummary(src_ip, dst_ip,
src_port, dst_port, tls_version))

    # Let's see how we did!
    print(f'There were {len(tls_connection_list)} connections
observed')
    for conn in tls_connection_list:
        if conn.get_tls_version() in ("SSL 3.0", "TLS 1.0", "TLS
1.1", "Unknown"):
            conn.print_summary()
```

```
[ ]    tls_version_audit('./sample.pcap')
```